

# Customer Experience Text Analytics for Social Media

Christoph Emunds, Benedikt Heinrichs, Dominik Nerger, Richard Polzin

March 16, 2017

## 1 Introduction

The project presented in this document focusses on finding and analyzing customers' opinions on products and their aspects from social media posts. This project plan describes the required background knowledge and the related work that we are going to build on. We present the current state of the project and explain what has been done so far. Finally, we define the goals for our project and give a schedule for the upcoming tasks.

## 2 Related work

### 2.1 Word embeddings

Word embeddings are numeric representations of words where similar words have similar vectors. Today, word embeddings are used as input to a lot of natural language processing (NLP) tasks, which is the reason why the process of training these vectors is sometimes called pre-training.

Word embeddings are generally more concise and have a lower dimensionality than the usual *bag of words* model with one-hot encoded vectors. Depending on the amount of words in the vocabulary, these one-hot vectors have a huge dimensionality, while being sparse (i.e. only one component of each vector is set to 1 and the others to 0) and wasting a lot of space. Such simple models also fail to incorporate meaning of and similarities between words. For example in the sentences

the cat got squashed in the garden on friday

the dog got flattened in the yard on monday

simple language models do not capture the similarities between the word pairs *cat/dog*, *squashed/flattened*, *garden/yard* and *friday/monday*.

A famous example for the expressive power of word embeddings is the following:

$$king - man + woman \approx queen$$

This equation shows that if one were to subtract the vector representation for the word *man* from the vector for *king* and add the vector for *woman*, this would result in a point near to the vector for the word *queen*.

The high dimensional vector representations of the words can be processed with the *t-SNE* dimensionality reduction algorithm to be able to plot them in a two dimensional vector space. Figure 1 shows a section of the vector space.




Figure 1: Representation of words relating to the differentiation between man and woman

## 2.2 POS-Tagging

POS-Tagging (part-of-speech tagging), also called grammatical tagging or word-category disambiguation, refers to the process of identifying partial parts of speech like nouns or verbs.

Identifying the role of a certain word within a sentence is important for the task of sentiment analysis, as identifying entities and their corresponding aspects can be handled much easier relying on certain assumptions about the part of speech of words.

Frequent nouns or noun phrases often describe aspects of products and the vocabulary that is used to describe those usually converges. With such an approach many important aspects can easily be found, but on a more general scheme grammatic based relationships can be used to extract important aspects.

For example words that express opinion on something, like 'great' or 'bad' can be incorporated in rules to extract aspects. To sum it up POS-Tagging is an important part of sentiment analysis.

While being important POS-Tagging is also a complex problem. Word-forms in natural language are often ambiguous. For example the word 'dogs' is usually thought of as a plural noun, but can be used as a verb as well:

The sailor dogs the hatch.

Due to the complexity, machine learning techniques are often applied in POS-Taggers. Popular approaches such as the Viterbi algorithm, the Brill tagger or the Baum-Welch algorithm work with techniques such as dynamic programming, supervised learning or hidden markov models.

## 2.3 NER-Tagging

Named-Entity recognition (NER) is a task that seeks to locate and classify specific information in text. This information is called a named entity and can refer to categories such as the names of

persons, locations, times or many others.

An annotated sentence could look like this :

[Tim Cook]<sub>[Person]</sub> has a Net worth of [785 million USD]<sub>[MonetaryValue]</sub> as of [March 16. 2017]<sub>[Time]</sub>.

Often the task of NER-Tagging is separated into two. In the first part names are detected and in the second part these names are classified.

While state-of-the-art NER Taggers perform very well and produce near-human performance they are also brittle and do not perform well in domains they were not designed for.

## 2.4 Dependency parsing

Dependency parsing is a method of building up context in a sentence. It looks on the relations each word has with the other words and tries to create a meaning out of a sentence by that.

For example the sentence "I saw a girl with a telescope" has two times the combination of an object with the word "a". The word "girl" is specified by the verb "saw" and the object "telescope" is specified by the word "with" with references to the verb "saw". The verb "saw" relates then to the noun "I" as it is the subject. A visualization can be seen in the following. This is however just one method of analyzing this sentence. Another way of seeing this sentence would be to relate the word "with" with the object "girl" which is possible in the language and gives the whole sentence a different meaning.

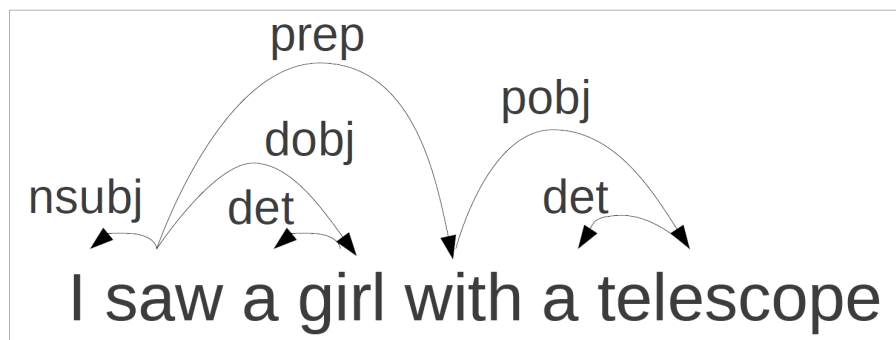


Figure 2: The dependency parsing example

This example has been taken from (Neubig, 2014).

## 3 Extracting entities and aspects

To be able to determine the customers' opinions on certain entities and their aspects, we first need to extract possible entities and aspects from the given dataset. For this task, we will use a method proposed by Hu and Liu (2004), which consists of the following two steps:

First, we identify frequent nouns and noun phrases with a POS-Tagger. Since people commenting on aspects of a product usually use a limited vocabulary, this process should yield a first set of possible entities and aspects. The frequency threshold needs to be chosen experimentally.

However, since the first step can miss quite a lot of important aspects, the next step tries to find some of them by exploiting the relationships between aspects and opinion words.

## 4 Determining sentiment polarities

For determining the sentiment polarities, we will focus on two approaches. The first one will be a supervised method, while the second is an unsupervised approach.

### 4.1 LSTM-Network

A recurrent neural network with long short-term memory (LSTM) cells is able to build a model over sequential data, which can be used for prediction. We will build an LSTM-network similar to the one of Hong and Fang (2015). However, training a recurrent neural network is a supervised task that requires some labeled input. Therefore, we need to apply a bootstrapping process first, which receives a collection of unlabelled sentences and a set of emotions and annotates the unlabelled sentences with the emotions expressed in each sentence. This should give us a labeled dataset of sufficient size to train the LSTM-network.

### 4.2 Lexicon-based approach

lexicon-based approach, because it is unsupervised and generally more powerful than supervised learning(?)

We will apply a method by Ding, Liu and Yu (2008), which has the following steps: First, all sentiment words and phrases in the text will be marked with the help of a sentiment lexicon.

Second, sentiment shifters will be applied.

Third, but-clauses must be handled.

Fourth, opinions need to be aggregated.

This method can be made more effective by applying dependency parsing to the sentences, such that the scope of each individual sentiment word can be determined more accurately. This allows us to discover the sentiment orientation of context dependent words (Bing Liu, 2012).

## 5 State of the project

So far, we built and tested the syntax parser *Parsey McParseface* (aka SyntaxNet) by Google, which we will use as POS-Tagger for extracting entities and aspects and as dependency parser for several other tasks. SyntaxNet is easy to use and yields state-of-the art results. Moreover, it is based on Google's Deep Learning Framework TensorFlow and therefore integrates nicely into the rest of our implementation.

We preprocessed the data by extracting the actual text of the posts into separate files, which can be processed by SyntaxNet.

Table 2 shows a comparison of several dependency parser models in terms of per-token accuracy.

We will use pre-trained word embeddings from the GloVe project.

We will use the sentiment lexicon by Bing Liu, which includes mis-spellings, morphological variants, slang and social-media mark-up of 2006 positive and 4783 negative words.

Table 1: Per-token accuracy of different POS-Taggers

Model	News	Web	Questions
Ling et al. (2015)	97.44	94.03	96.18
Andor et al. (2016)	97.77	94.80	96.86
Parsey McParseface	97.52	94.24	96.45

Table 2: Accuracy of different dependency parsers

Model	News	Web	Questions
Martins et al. (2013)	93.10	88.23	94.21
Zhang and McDonald (2014)	93.32	88.65	93.37
Weiss et al. (2015)	93.91	89.29	94.17
Andor et al. (2016)	94.44	90.17	95.40
Parsey McParseface	94.15	89.98	94.77

## 6 Goals and schedule

The overall goal is to extract as many quintuples  $(e_i, a_{ij}, s_{ijkl}, h_k, t_l)$  as possible, where  $e_i$  is the  $i$ 'th entity and  $a_{ij}$  is the  $j$ 'th aspect of entity  $i$  the opinion is expressed on.  $s_{ijkl}$  is the sentiment polarity, which can take on the values *positive* or *negative*.  $h_k$  describes the opinion holder and  $t_l$  the time at which the opinion was expressed.

We aim to build a database with opinions and statistics on several products and their aspects. For a visualization of the results we will build a simple dashboard with state-of-the-art web technologies.

## References

- Andor, D., Alberti, C., Weiss, D., Severyn, A., Presta, A., Ganchev, K., ... Collins, M. (2016). Globally normalized transition-based neural networks. *CoRR*, *abs/1603.06042*. Retrieved from <http://arxiv.org/abs/1603.06042>
- Ding, X., Liu, B., & Yu, P. S. (2008). A holistic lexicon-based approach to opinion mining. In *Proceedings of the 2008 international conference on web search and data mining* (pp. 231–240). New York, NY, USA: ACM. Retrieved from <http://doi.acm.org/10.1145/1341531.1341561> doi: 10.1145/1341531.1341561
- Hong, J., & Fang, M. (2015). Sentiment analysis with deeply learned distributed representations of variable length texts.
- Hu, M., & Liu, B. (2004). Mining and summarizing customer reviews. In *Proceedings of the tenth acm sigkdd international conference on knowledge discovery and data mining* (pp. 168–177). New York, NY, USA: ACM. Retrieved from <http://doi.acm.org/10.1145/1014052.1014073> doi: 10.1145/1014052.1014073
- Liu, B. (2012). *Sentiment analysis and opinion mining*.
- Neubig, G. (2014). *Nlp programming tutorial 12 - dependency parsing*. Retrieved from [www.phontron.com/slides/nlp-programming-en-11-depend.pdf](http://www.phontron.com/slides/nlp-programming-en-11-depend.pdf)
- Pennington, J., Socher, R., & Manning, C. D. (2014). Glove: Global vectors for word representation. In *Empirical methods in natural language processing (emnlp)* (pp. 1532–1543). Retrieved from <http://www.aclweb.org/anthology/D14-1162>
- Presta, A., Severyn, A., Golding, A., Bohnet, B., Alberti, C., Andor, D., ... Credo, T. (n.d.). *Syntaxnet: Neural models of syntax*. Retrieved from <https://github.com/tensorflow/models/tree/master/syntaxnet>
- Zhang, L., & Liu, B. (2014). Aspect and entity extraction for opinion mining. In W. W. Chu (Ed.), *Data mining and knowledge discovery for big data: Methodologies, challenge and opportunities* (pp. 1–40). Berlin, Heidelberg: Springer Berlin Heidelberg. Retrieved from [http://dx.doi.org/10.1007/978-3-642-40837-3\\_1](http://dx.doi.org/10.1007/978-3-642-40837-3_1) doi: 10.1007/978-3-642-40837-3\_1